

C# Programming

Dr.-Ing. Grigory Devadze

September 5, 2025

* :BEAMER_{env}: frame :BEAMER_{envargs}: [plain]

Interfaces

Why Interfaces?

- Define **contracts** between classes
- Separate **what** is done from **how** it is done
- Enable testability and interchangeability of components
- Foundation for Dependency Injection and loose coupling

Interface Syntax in C#

```
public interface ILogger
{
    void Log(string message);
}
```

Implementation

```
public class ConsoleLogger : ILogger
{
    public void Log(string message) =>
        Console.WriteLine(message);
}
```

Example Program

```
public static class ProgramInterfaceDemo
{
    public static void Main()
    {
        ILogger log = new ConsoleLogger();
        log.Log("Hello from Interface demo!");
    }
}
```

Multiple Inheritance

Example

```
public interface IShape { double Area(); }
public interface IColored { string Color { get; } }

public class Circle : IShape, IColored
{
    public double Radius { get; set; }
    public string Color { get; set; } = "red";
    public double Area() =>
        Math.PI * Radius * Radius;
}
```

Typical Use Cases

- Logging (ILogger → FileLogger, ConsoleLogger, RemoteLogger)
- Repositories (IDataRepository → SqlRepository, InMemoryRepository)
- UI/Events (IObserver / IObservable)
- Strategy pattern and plug-in systems

Why the “I” in ILogger?

- In C#, interface names are conventionally prefixed with I
 - Examples: ILogger, IDisposable, IEnumerable
- Origin: COM (Component Object Model) and early .NET guidelines
- Helps distinguish interfaces from classes at a glance
- Not enforced by the compiler → you **could** write “Logger”
- But: community and framework standards expect “I” prefix

Loose Coupling

What is Coupling?

- Coupling = how tightly two classes depend on each other
- **Tight coupling**: a class directly uses another concrete class
- **Loose coupling**: a class only depends on an **interface**, not a specific implementation

Example: Tight Coupling ()

```
public class ReportService
{
    private readonly ConsoleLogger _logger = new ConsoleLogger();

    public void Generate()
    {
        _logger.Log("Generating report...");
    }
}
```

Loose Coupling

Example: Loose Coupling ()

```
public class ReportService
{
    private readonly ILogger _logger;
    public ReportService(ILogger logger) => _logger = logger;

    public void Generate()
    {
        _logger.Log("Generating report...");
    }
}
```

Benefits

- Flexible: swap ConsoleLogger → FileLogger → RemoteLogger
- Easier to maintain and extend
- Foundation for Dependency Injection

What is Dependency Injection (DI)?

- **Dependency** = a required component (e.g., ILogger)
- **Injection** = passing the dependency **from outside** instead of creating it inside
- Promotes **loose coupling** and **testability**
- Widely used with DI frameworks (e.g., Microsoft.Extensions.DependencyInjection, Autofac)

Who depends on whom?

- Provider: implements the interface (e.g., ConsoleLogger : ILogger)
- Consumer: **uses** the interface (e.g., ReportService needs ILogger)
- Goal: Consumers depend on **abstractions**, not concrete classes

Dependency Injection

Example: Constructor Injection

```
public class ReportService
{
    private readonly ILogger _logger;

    // Dependency injected here
    public ReportService(ILogger logger)
    {
        _logger = logger;
    }

    public void Generate()
    {
        _logger.Log("Generating report...");
    }
}
```

Dependency Injection

// Usage

```
ILogger log = new ConsoleLogger();  
var service = new ReportService(log);  
service.Generate();
```

Benefits

- Easily swap implementations (ConsoleLogger → FileLogger → RemoteLogger)
- Enables unit testing with **mock** loggers
- Forms the foundation of modern .NET applications

Conceptual Differences

- **Interface** = contract (what must be done)
- **Abstract class** = contract + shared implementation + state
- Inheritance:
 - Interfaces → multiple allowed
 - Abstract class → single base class only
- State/constructors:
 - Interfaces → no instance fields
 - Abstract classes → can have fields, constructors, protected helpers

Syntax Overview

```
public interface IWorker
{
    void DoWork();
    // since C# 8: default method implementations allowed
    void Report() => Console.WriteLine("Work done.");
}

public abstract class WorkerBase
{
    protected readonly string Name;
    protected WorkerBase(string name) => Name = name;

    public abstract void DoWork();           // must override
    public virtual void Report() =>         // may override
        Console.WriteLine($"{Name}: done.");
}
```

When to Choose Which?

- Choose **interface** when:
 - You need a **contract** only; implementations vary widely
 - You want consumers to be loosely coupled (DI)
 - Multiple, orthogonal capabilities are needed (e.g., ILogger, IAsyncDisposable)
- Choose **abstract class** when:
 - Implementations share **state** and **non-trivial common code**
 - You need **protected** helpers, **constructors**, or **template methods**
 - You want to evolve behavior in a single place

Advanced: Explicit Interface Implementation

Why?

- To avoid name clashes (two interfaces with same method name)
- To **hide** interface members from the public API

```
public interface IPrintable { void Print(); }  
public class Report : IPrintable  
{  
    void IPrintable.Print() =>  
        Console.WriteLine("Report printed");  
}
```

```
var r = new Report();  
// r.Print(); // not visible  
(IPrintable)r.Print(); // works
```

Heavily used in .NET Core / ASP.NET Core

- `ILogger` → Logging abstraction
- `IConfiguration` → App settings
- `IServiceCollection` → Dependency Injection
- `IEnumerable<T>` → Collections and LINQ
- `IAsyncDisposable` → Async cleanup

Takeaway

- Most modern .NET features are built on interfaces
- Learning to design with interfaces = understanding the .NET ecosystem

Combining independent behaviors

```
public interface IPrintable { void Print(); }  
public interface ISavable  { void Save(string path); }  
  
public class Report : IPrintable, ISavable  
{  
    public void Print() => Console.WriteLine("Printing...");  
    public void Save(string p) => File.WriteAllText(p, "data");  
}
```

- Interfaces let you compose independent features without deep inheritance

“I-itis”

- Defining interfaces just for the sake of it (IDoSomething)
- Abstractions without a clear purpose add complexity

Over-abstraction

- Don't create interfaces until you **need multiple implementations**
- Prefer YAGNI (You Aren't Gonna Need It)
- Abstract class is safer if API must evolve often

* :BEAMER_{env}: frame :BEAMER_{envargs}: [plain]

Events

Why Events?

- Mechanism for publish/subscribe communication
- Enables **loose coupling** between components
- Core to many .NET frameworks (UI, I/O, async notifications)

```
public class Publisher
{
    public event EventHandler? SomethingHappened;

    public void DoWork()
    {
        Console.WriteLine("Working...");
        SomethingHappened?.Invoke(this, EventArgs.Empty);
    }
}

public class Subscriber
{
    public void OnSomething(object? sender, EventArgs e) =>
        Console.WriteLine("Received event!");
}
```


Example Program

```
public static class ProgramEventDemo
{
    public static void Main()
    {
        var pub = new Publisher();
        var sub = new Subscriber();

        pub.SomethingHappened += sub.OnSomething;

        pub.DoWork(); // Triggers event
    }
}
```

Passing additional data

```
public class DataEventArgs : EventArgs
{
    public string Message { get; }
    public DataEventArgs(string msg) => Message = msg;
}

public class Publisher
{
    public event EventHandler<DataEventArgs>? DataAvailable;

    public void Publish(string msg) =>
        DataAvailable?.Invoke(this, new DataEventArgs(msg));
}
```

FileSystemWatcher

```
using System.IO;

var watcher = new FileSystemWatcher(".");
watcher.Created += (s,e) => Console.WriteLine($"File created: {e.Name}");
watcher.EnableRaisingEvents = true;

Console.WriteLine("Press Enter...");
Console.ReadLine();
```

Timer

```
using var timer = new System.Timers.Timer(1000); // 1s
timer.Elapsed += (s,e) => Console.WriteLine("Tick");
timer.Start();

Console.ReadLine();
```

Anonymous Event Handlers

```
public class ReportService
{
    private readonly ILogger _logger;

    // Dependency injected here
    public ReportService(ILogger logger)
    {
        _logger = logger;
    }

    public void Generate()
    {
        _logger.Log("Generating report...");
    }
}
```

Unsubscribing

Multicast Delegates

- An event can notify multiple subscribers
- Subscribers are invoked in order of subscription
- All receive the same event arguments

```
pub.SomethingHappened += (s,e) => Console.WriteLine("Handler 1");  
pub.SomethingHappened += (s,e) => Console.WriteLine("Handler 2");  
pub.DoWork();
```

Output

Working...

Handler 1

Handler 2

Events in C#

- Safe, built-in publish/subscribe pattern
- Provide loose coupling between components
- Can have multiple subscribers (multicast)
- Used throughout .NET (UI, I/O, diagnostics, async notifications)
- Always unsubscribe in long-running apps to avoid memory leaks

Asynchronous Programming

Why Asynchronous Programming?

- Improve responsiveness (UI, servers stay reactive)
- Scale with I/O-bound work (network, disk, database)
- Avoid blocking threads unnecessarily
- Enable modern patterns: streaming or microservices

Sync vs Async vs Parallel

- **Synchronous**: one thing at a time, blocking
- **Asynchronous**: start an operation, continue without waiting
- **Parallel**: multiple threads doing work **at the same time**
- Async \neq Parallel!

Tasks

- Task represents an ongoing operation
- Task<T> represents an operation that returns a value
- Tasks can be awaited, combined, or cancelled

Async/Await Keywords

```
public async Task<int> ComputeAsync()  
{  
    await Task.Delay(1000); // simulate async work  
    return 42;  
}
```

```
public static class ProgramAsyncDemo
{
    public static async Task Main()
    {
        Console.WriteLine("Starting...");
        int result = await ComputeAsync();
        Console.WriteLine($"Result = {result}");
    }

    static async Task<int> ComputeAsync()
    {
        await Task.Delay(1000);
        return 0;
    }
}
```

File I/O

```
string text = await File.ReadAllTextAsync("data.txt");  
Console.WriteLine($"Read {text.Length} chars");
```

HTTP Requests

```
using var http = new HttpClient();  
string html = await http.GetStringAsync("https://example.com");  
Console.WriteLine(html.Substring(0,100));
```

WhenAll and WhenAny

```
var t1 = http.GetStringAsync("https://example.com");  
var t2 = http.GetStringAsync("https://dotnet.microsoft.com");  
  
await Task.WhenAll(t1, t2);    // wait for both  
Console.WriteLine(t1.Result.Length + t2.Result.Length);  
  
var first = await Task.WhenAny(t1, t2); // first to finish  
Console.WriteLine("First completed!");
```

Exceptions in Async

```
try
{
    await DangerousAsync();
}
catch(Exception ex)
{
    Console.WriteLine($"Caught: {ex.Message}");
}
```

AggregateException

- Multiple exceptions can occur (e.g., Task.WhenAll)
- Flattened automatically when awaiting
- Example: await Task.WhenAll(tasks) throws the first exception

CancellationToken

```
var cts = new CancellationTokenSource();

var task = Task.Run(async () =>
{
    while (true)
    {
        cts.Token.ThrowIfCancellationRequested();
        await Task.Delay(500);
    }
}, cts.Token);

// cancel after 2s
cts.CancelAfter(2000);
try { await task; }
catch (OperationCanceledException) { Console.WriteLine("Cancelled!"); }
```

Timeouts and Throttling

Timeout

```
using var cts = new CancellationTokenSource(TimeSpan.FromSeconds(3));  
await http.GetStringAsync("https://slow.example", cts.Token);
```

Throttling Parallelism

```
var urls = new [] { "a", "b", "c" };  
using var sem = new SemaphoreSlim(3);  
var tasks = urls.Select(async url =>  
{  
    await sem.WaitAsync();  
    try  
    {  
        var data = await http.GetStringAsync(url);  
        Console.WriteLine($"{url}: {data.Length}");  
    }  
    finally { sem.Release(); }  
});  
await Task.WhenAll(tasks);
```


`IProgress<T>`

```
public async Task DownloadAsync(IProgress<int> progress)
{
    for(int i=0; i<=100; i+=10)
    {
        await Task.Delay(100);
        progress.Report(i);
    }
}
```

```
var p = new Progress<int>(v => Console.WriteLine($"{v}%"));
await DownloadAsync(p);
```

Async Streams

```
IAsyncEnumerable<T>
```

```
public async IAsyncEnumerable<int> NumbersAsync()  
{  
    for (int i=0; i<5; i++)  
    {  
        await Task.Delay(200);  
        yield return i;  
    }  
}
```

```
await foreach(var n in NumbersAsync())  
    Console.WriteLine(n);
```

Pitfalls to Avoid

- Blocking async code with `.Result` or `.Wait()` → deadlocks!
- Forgetting `await` → tasks run but exceptions disappear
- Mixing sync + async badly → thread pool starvation
- Ignoring `CancellationToken`s

Guidelines

- Use `async/await` all the way down
- Prefer `ConfigureAwait(false)` in libraries
- Always support cancellation if possible
- Keep `async` methods non-blocking

Asynchronous Programming in C#

- Async/await simplifies non-blocking code
- Tasks model ongoing work, can be awaited, combined, cancelled
- Great for I/O-bound scalability
- Not the same as threading
- Foundation for reactive and streaming systems

Threading

Why Threading?

- Exploit multiple CPU cores for concurrency
- Run multiple operations simultaneously
- Keep applications responsive (background work)
- Essential for parallel computing and servers

Async vs Threading

- **Asynchronous**: efficient I/O, non-blocking
- **Threading**: true parallelism for CPU-bound work
- Often combined in modern apps

Creating Threads

```
using System.Threading;

Thread t = new Thread(() =>
{
    Console.WriteLine("Hello from worker thread!");
});
t.Start();
t.Join(); // wait for completion
```

Output

Hello from worker thread!

ThreadPool and Tasks

- Creating raw threads is expensive
- .NET provides a **ThreadPool** for short-lived tasks
- Usually accessed via `Task.Run`

```
Task t = Task.Run(() =>
{
    Console.WriteLine("Running on ThreadPool");
});
t.Wait();
```

Race Conditions

```
int counter = 0;

Parallel.For(0, 10000, i =>
{
    counter++; // not thread-safe!
});
Console.WriteLine(counter); // not 10000
```

Why? Multiple threads update the same memory concurrently

Lock

```
object locker = new();  
int counter = 0;
```

```
Parallel.For(0, 10000, i =>  
{  
    lock(locker)  
    {  
        counter++;  
    }  
});  
Console.WriteLine(counter); // always 10000
```

Interlocked

```
int counter = 0;

Parallel.For(0, 10000, i =>
{
    Interlocked.Increment(ref counter);
});
Console.WriteLine(counter); // always 10000
```

Lock vs Interlocked vs No Sync

- Example results (N=5,000,000, 12 cores)
 - lock: a=5000000, time=156 ms
 - Interlocked: b=5000000, time=124 ms
 - NO sync (!!): c=417251, time=38 ms

Insights

- Correctness > performance
- Interlocked is faster for simple counters
- Locks allow more complex critical sections

Many readers, few writers

```
var dict = new Dictionary<int,int>();
var rw = new ReaderWriterLockSlim();

Parallel.For(0, 10000, i =>
{
    if (i % 10 == 0)
    {
        rw.EnterWriteLock();
        try { dict[i] = i; }
        finally { rw.ExitWriteLock(); }
    }else
    {
        rw.EnterReadLock();
        try { _ = dict.ContainsKey(i - 1); }
        finally { rw.ExitReadLock(); }
    }
});
```

Thread-Local Storage

ThreadLocal<T>

```
var local = new ThreadLocal<int>(() => 0);
```

```
Parallel.For(0, 1000, i =>
{
    local.Value++;
    Console.WriteLine($"Thread \
                        {Thread.CurrentThread.ManagedThreadId}: \
                        {local.Value}");
});
```

Each thread maintains its own independent value

Common Issues

- Race conditions (unprotected shared state)
- Deadlocks (threads waiting forever)
- Starvation (some threads never scheduled)
- Excessive threads → overhead, context switching

Guidelines

- Prefer `Task.Run` and the `ThreadPool` for most work
- Use `lock` for critical sections, `Interlocked` for counters
- Use `ReaderWriterLockSlim` for read-heavy scenarios
- Avoid creating too many raw threads
- Combine with `async` for modern scalable systems

Threading in C#

- Enables true CPU parallelism
- Threads are heavy, prefer the ThreadPool
- Synchronization is essential for correctness
- Always balance correctness vs performance
- Foundation for higher-level abstractions (TPL, Parallel LINQ, async/await)

Diagnostics

Why Diagnostics?

- Production systems need visibility:
 - Debugging errors
 - Performance bottlenecks
 - Monitoring live behavior
- `Console.WriteLine` is not enough
- Structured logging and events enable tools and automation

Key Concepts

- **Logging** = what happened
- **Tracing** = sequence of events
- **Metrics** = numeric measurements over time

System.Diagnostics

```
using System.Diagnostics;  
  
Trace.Listeners.Add(new TextWriterTraceListener("log.txt"));  
  
Trace.WriteLine("Application started");  
Trace.TraceWarning("Low memory");  
Trace.TraceError("Something went wrong");  
Trace.Flush();
```

Notes

- Works with multiple listeners (console, file, custom)
- Simple, but limited compared to modern logging

Structured Events (ETW-style)

```
using System.Diagnostics.Tracing;

[EventSource(Name = "MyApp-Events")]
public sealed class MyEventSource : EventSource
{
    public static readonly MyEventSource Log = new();

    [Event(1, Level=EventLevel.Informational)]
    public void RequestStart(string route) => WriteEvent(1, route);

    [Event(2, Level=EventLevel.Informational)]
    public void RequestStop(int status) => WriteEvent(2, status);

    [Event(3, Level=EventLevel.Error)]
    public void Failure(string message) => WriteEvent(3, message);
}
```

Usage

```
MyEventSource.Log.RequestStart("/api");  
MyEventSource.Log.Failure("DB connection failed");
```

Tools

- Capture with dotnet-trace collect --providers MyApp-Events
- View in PerfView or Visual Studio diagnostics

ILogger in .NET Core

```
using Microsoft.Extensions.Logging;  
  
var factory = LoggerFactory.Create(b => b.AddConsole());  
ILogger log = factory.CreateLogger("Demo");  
  
log.LogInformation("Started processing");  
log.LogWarning("Something looks odd");  
log.LogError("An error occurred");
```

Benefits

- Unified abstraction
- Multiple providers (console, file, cloud, Application Insights)
- Configurable log levels

EventCounter API

```
public sealed class MetricsSource : EventSource
{
    public static readonly MetricsSource Log = new("MetricsDemo");

    private EventCounter _reqPerSec;

    public MetricsSource(string name) : base(name)
    {
        _reqPerSec = new EventCounter("requests-per-second", this);
    }

    public void Report(double value) => _reqPerSec.WriteMetric(value);
}
```

Tools

- `dotnet-counters` monitor `MetricsDemo`
- Real-time metrics: request rate, latency, memory usage

Low-level Instrumentation

```
using System.Diagnostics;

var source = new DiagnosticListener("MyLibrary");

if (source.IsEnabled("RequestStart"))
    source.Write("RequestStart", new { Url = "/api", Method = "GET" });
```

Observers

- Libraries emit diagnostic events
- Observability systems (OpenTelemetry, Application Insights) subscribe
- Decouples producers and consumers

.NET CLI Tools

- `dotnet-trace` → collect EventSource + runtime events
- `dotnet-counters` → monitor performance counters live
- `dotnet-dump` → capture and analyze process memory dumps
- PerfView → powerful UI for ETW / EventPipe traces

Example

```
dotnet-trace collect --process-id 1234 --providers MyApp-Events  
dotnet-counters monitor System.Runtime
```

Diagnostics and Instrumentation

- Logging, tracing, metrics = 3 pillars of observability
- Use ILogger for app-level logging
- Use EventSource and EventCounter for structured diagnostics
- Use DiagnosticSource for library instrumentation
- .NET tools (trace, counters, dump, PerfView) provide visibility into live systems

Reflection

Why Reflection?

- Inspect types, methods, properties at runtime
- Dynamically invoke members
- Discover metadata (attributes, interfaces)

Why Attributes?

- Declaratively attach metadata to code
- Frameworks use attributes to drive behavior
- Cleaner than hardcoding configuration in code

Get Type Information

```
Type t1 = typeof(string);  
Type t2 = "hello".GetType();
```

```
Console.WriteLine(t1.FullName); // System.String  
Console.WriteLine(t2.IsClass);  // True
```

Inspect Members

```
Type t = typeof(DateTime);  
  
foreach(var m in t.GetMethods())  
    Console.WriteLine(m.Name);
```


Invoke a Method

```
Type t = typeof(Math);  
object? result = t.InvokeMember("Sqrt",  
    BindingFlags.InvokeMethod  
    | BindingFlags.Public  
    | BindingFlags.Static,  
    binder: null,  
    target: null,  
    args: new object[] { 9.0 });  
  
Console.WriteLine(result); // 3
```

Create Instance Dynamically

```
Type t = typeof(StringBuilder);  
object obj = Activator.CreateInstance(t);
```

Load Types from Assembly

```
var asm = Assembly.GetExecutingAssembly();  
foreach (var type in asm.GetTypes())  
    Console.WriteLine(type.FullName);
```

Use Cases

- Plugin systems (discover types at runtime)
- Reflection-based DI containers
- Unit test discovery

Common Examples

```
[Obsolete("Use NewMethod instead")]
public void OldMethod() {}

[Serializable]
public class Person { }

public class Customer
{
    [Required]
    public string Name { get; set; }
}
```

Define Your Own

```
[AttributeUsage(AttributeTargets.Class | AttributeTargets.Method)]  
public class AuthorAttribute : Attribute  
{  
    public string Name { get; }  
    public AuthorAttribute(string name) => Name = name;  
}
```

Apply It

```
[Author("Me")]  
public class Report { }
```

With Reflection

```
Type t = typeof<Report>;  
var attrs = t.GetCustomAttributes(typeof<AuthorAttribute>,  
                                   inherit: false);  
  
foreach (AuthorAttribute a in attrs)  
    Console.WriteLine(a.Name); // Me
```

Attributes in Frameworks

- ASP.NET Core routing:
 - `[HttpGet("/api/data")]`
- Entity Framework:
 - `[Key]`, `[Table("Users")]`
- NUnit/xUnit testing:
 - `[Test]`, `[Fact]`
- Serialization:
 - `[JsonProperty("id")]`

Benefits

- Declarative, less boilerplate
- Frameworks discover behavior via reflection
- Extensible → you can define your own

Reflection

- Slower than direct code (metadata lookup)
- Can bypass encapsulation (dangerous if misused)
- Complicates code analysis and refactoring

Attributes

- Overuse can clutter code
- Hard to validate at compile time
- Framework coupling (depends on specific attributes)

Reflection and Attributes

- Reflection = runtime type inspection & dynamic invocation
- Attributes = declarative metadata
- Together they power many frameworks (ORMs, DI, test frameworks, serializers)
- Use reflection with care: correctness first, performance second
- Attributes are a clean way to configure behavior declaratively

Native Interop

Why Native Interop?

- Reuse existing C/C++ libraries
- Access operating system APIs directly (Win32, libc)
- Performance: use optimized native code (e.g., BLAS, OpenSSL)
- Device integration (USB, PCIe, hardware APIs)

Approaches

- P/Invoke (Platform Invocation Services)
- COM Interop (legacy Windows components)
- C++/CLI bridge (mixed-mode assemblies)
- Focus: **P/Invoke** = most common & practical

Importing Native Functions

```
using System.Runtime.InteropServices;

public static class NativeMethods
{
    [DllImport("user32.dll", CharSet = CharSet.Unicode)]
    public static extern int MessageBox(
        IntPtr hWnd,
        string text,
        string caption,
        uint type);
}
```

Usage

```
NativeMethods.MessageBox(IntPtr.Zero,
    "Hello from native code!",
    "Interop Demo", 0);
```

Automatic Type Conversion

- P/Invoke marshals managed \leftrightarrow unmanaged types
- Common cases:
 - `string` \rightarrow `LPWSTR` / `char*`
 - `bool` \rightarrow `BOOL`
 - `int`, `double` \rightarrow `int`, `double`
- Attributes customize marshaling:
 - `[MarshalAs(UnmanagedType.LPStr)] string s;`

Structs

```
dotnet-trace collect --process-id 1234 --providers MyApp-Events  
dotnet-counters monitor System.Runtime
```

Calling Native Code with P/Invoke

```
// Windows API (C)
DWORD GetTickCount(void);

// C# P/Invoke declaration
using System.Runtime.InteropServices;

public static class NativeMethods
{
    [DllImport("kernel32.dll")]
    public static extern uint GetTickCount();
}

Console.WriteLine(
    $"Ticks: {NativeMethods.GetTickCount()}");
);
```

Notes

- 'DWORD' = 32-bit unsigned integer ('uint' in C#).
- Function returns system uptime in ms (wraps after ~49.7 days).
- Shows how WinAPI types map to C# types in interop.

Passing Delegates to Native Code

```
[UnmanagedFunctionPointer(CallingConvention.Cdecl)]  
public delegate int Callback(int x);  
  
[DllImport("NativeLib.dll")]  
public static extern void RegisterCallback(Callback cb);  
  
// Usage  
RegisterCallback(x => {  
    Console.WriteLine($"Callback {x}");  
    return x * 2;  
});
```

Marshaling delegates allows native → managed calls

Unsafe Code

```
unsafe
{
    int value = 42;
    int* ptr = &value;
    Console.WriteLine(*ptr); // dereference
}
```

Notes

- Needed for high-performance pointer manipulation
- Use sparingly in safe code bases

Where P/Invoke is Used

- Call Windows APIs (MessageBox, file operations)
- Bind to C libraries (e.g., SQLite, OpenSSL, LAPACK)
- Interact with hardware drivers
- Performance hotspots (SIMD, image processing)

Interop Risks

- Crashes (bad pointers, invalid structs)
- Memory leaks (unmanaged allocations not GC'd)
- Performance overhead from marshaling
- Platform differences (x86/x64, Windows/Linux)

Guidelines

- Prefer managed APIs when available
- Isolate interop in dedicated classes (NativeMethods)
- Always match signatures exactly (calling convention, struct layout)
- Test on all target platforms
- Use safe handles for unmanaged resources

Native Interop in C#

- P/Invoke allows direct calls to native DLLs
- Marshaling bridges managed and unmanaged worlds
- Useful for OS APIs, performance, device integration
- Powerful but dangerous: handle with care
- Foundation for system programming with .NET

Garbage collection

Why care about memory & GC?

- Throughput & latency depend on allocation patterns
- Latency-sensitive workloads (UI, trading, games) need predictable pauses
- Cloud efficiency: fewer GB → fewer €€
- Debuggability: find leaks, finalizer issues

Value vs Reference types

- Value types (struct) usually live on the stack
- Reference types (class) live on the managed heap
- Boxing allocates: `object o = 1;` → avoid in hot paths

Stack vs Heap

- **Stack**: fast, scoped, no GC
- **Heap**: managed by GC, supports long-lived objects

Generational GC

- Gen0 → short-lived objects
- Survivors → Gen1 → Gen2 (long-lived)
- Background GC for Gen2; ephemeral segments for Gen0/Gen1

Heaps

- SOH (Small Object Heap): $< \sim 85$ KB
- LOH (Large Object Heap): $\geq \sim 85$ KB (collected with Gen2, prone to fragmentation)
- POH (Pinned Object Heap): isolates pinned blocks to reduce SOH/LOH fragmentation

Fast path & pitfalls

// Fast: small arrays/objects on SOH

```
var buf = new byte[1024];
```

// LOH allocation (>= ~85 KB) - more expensive & fragmented

```
var big = new byte[200_000];
```

- Many small, short-lived allocs → OK (Gen0 friendly)
- Many large allocs → LOH churn, consider pooling (reusing instead new T[n])

LOH churn

- Frequent alloc/free cycle = “LOH churn” → more Gen2 collections, fragmentation

Mitigation: pooling

- Reuse large buffers instead of allocating each time
- `ArrayPool<T>.Shared.Rent(size)` → get array (may be larger than requested)
- Use the array, then `ArrayPool<T>.Shared.Return(array)`
- Optionally clear on return (`Return(array, clearArray:true)`) for security

Example

```
using System.Buffers;

// Rent a 100k-element array from the shared pool
int[] buffer = ArrayPool<int>.Shared.Rent(100_000);

try {
    buffer[0] = 1;
}
finally {
    // IMPORTANT: return for reuse
    ArrayPool<int>.Shared.Return(buffer);
}
```

- Avoid holding rented arrays for long periods
- Great for networking, serialization, image/audio/video processing
- Cuts down LOH allocations → fewer pauses, less fragmentation

Collection counts

```
for (int gen = 0; gen <= GC.MaxGeneration; gen++)  
    Console.WriteLine($"Gen{gen}: {GC.CollectionCount(gen)}");
```

Total managed memory

```
Console.WriteLine($"Managed: \  
    {GC.GetTotalMemory(false)/1024.0:N1} KB");
```

For secrets (passwords, keys, tokens)

- Clear manually before releasing:
 - `Array.Clear(myArray)`
 - `CryptographicOperations.ZeroMemory(span)`
`(System.Security.Cryptography)`
- Prefer `SecureString` or `Span<byte>` buffers you control

Latency modes & NoGCRegion

```
GCSettings.LatencyMode = GCLatencyMode.SustainedLowLatency;  
// Critical section...
```

```
GCSettings.LatencyMode = GCLatencyMode.Interactive;
```

```
// Try NoGCRegion for short time-critical windows
```

```
if (GC.TryStartNoGCRegion(10_000_000))  
{ /* critical */ GC.EndNoGCRegion(); }
```

Modes

- **Workstation GC**: desktop apps, responsive UI
- **Server GC**: multi-core servers, multiple heaps, higher throughput
- Configure in runtimeconfig.json or hosting (ASP.NET Core defaults to Server GC)

Finalizers are non-deterministic

```
class Demo : IDisposable {  
    ~Demo() { /* non-deterministic, on finalizer thread */ }  
    public void Dispose() { /* deterministic */ GC.SuppressFinalize(this); }  
}
```

- Prefer IDisposable / using (or await using) over finalizers
- Use SafeHandle instead of raw IntPtr for native resources

Span<T> / Memory<T> (stackalloc, ref struct)

```
Span<int> stackData = stackalloc int[128];  
stackData[0] = 123;
```

```
// slicing without allocations
```

```
ReadOnlySpan<char> slice = "abcdef".AsSpan(1,3); // "bcd"
```

- Spans avoid heap allocations for slicing/parsing
- Great for serializers, parsing, IO pipelines

Reduce string churn

```
var sb = new System.Text.StringBuilder();  
for (int i=0; i<1000; i++) sb.Append(i).Append(',');  
string s = sb.ToString(); // allocate once
```

- Avoid “+” in loops; prefer StringBuilder
- Consider string.Create for formatted output in hot paths

State machines & ValueTask

```
public async ValueTask<int> FooAsync() { await Task.Yield(); return 1; }
```

- Each async method creates a state machine object **when it awaits**
- For frequently completing synchronously, consider ValueTask

Native interop needs stable pointers

Pinning hurts the GC

```
// Pinned memory can fragment the heap  
var handle = GCHandle.Alloc(buffer, GCHandleType.Pinned);  
// ...  
handle.Free();
```

- Prefer fixed `Span<byte>` patterns or POH; minimize pin duration

.NET CLI & tools

- `dotnet-counters monitor System.Runtime` (alloc rate, GC heap size, %time in GC)
- `dotnet-trace collect` (GC pauses, allocations)
- `dotnet-gcdump collect & analyze` in PerfView/VS
- Visual Studio: Diagnostic Tools (Memory Usage, Heap snapshots)
- PerfView: GC stats, LOH fragmentation, allocation stacks

Quick live check

```
dotnet-counters monitor --process-id <PID> System.Runtime  
# GC Heap Size, Gen0/1/2 collections, Allocation Rate, % Time in GC
```

BenchmarkDotNet (recommended)

```
[MemoryDiagnoser]
public class AllocBench {
    [Benchmark] public int Boxing() { object o = 42; return (int)o; }
    [Benchmark] public int NoBox() { int x = 42; return x; }
}
```

- Reports time + allocations (Gen0/LOH) per operation

Watch out for

- Loops with hidden allocations (boxing, LINQ closures, foreach on non-struct enumerators)
- Large temporary arrays → LOH churn
- Long-lived event subscriptions → leaks (publisher holds subscriber)
- Finalizer backlog → high GC pause times
- Excessive small string concatenations

Quick wins

- Reuse buffers (ArrayPool)
- Prefer structs for tiny immutable data on hot paths (avoid boxing)
- Avoid capturing lambdas in tight loops
- Consider TryParse over Parse (no exceptions → no allocations)
- Use logging templates (structured) instead of string interpolation in hot paths

Quick demos (copy-paste)

Gen counts & allocs

```
for (int i=0; i<10_000; i++) _ = new byte[1024];  
Console.WriteLine($"Gen0:{GC.CollectionCount(0)} Gen1:{GC.CollectionCount(1)}");
```

LOH impact

```
var sw = System.Diagnostics.Stopwatch.StartNew();  
for (int i=0; i<200; i++) _ = new byte[100_000]; // LOH  
sw.Stop();  
Console.WriteLine($"Time: {sw.ElapsedMilliseconds} ms");
```

GC settings

- Server vs Workstation GC (runtimeconfig.json, hosting)
- Heap hard limits in containers (respect cgroup memory)
- Environment variables: `COMPlus_GCHeapHardLimit`,
`DOTNET_GCHeapHardLimit`

Memory & GC in .NET

- Generational GC favors many short-lived objects
- LOH/POH require special care (pooling, pinning)
- Prefer `IDisposable`, pooled buffers, spans, and zero-alloc APIs
- Measure with `BenchmarkDotNet`; observe with `dotnet-counters` / `PerfView` / VS
- Tune only after measuring; correctness & clarity first

Why Serial Today?

- Industrial devices (PLCs, sensors, motor controllers)
- Embedded (MCUs, Arduino), lab instruments (SCPI)
- Simple, reliable, widely available (USB-to-Serial)

List Available Ports

```
using System;  
using System.IO.Ports;  
  
Console.WriteLine("Ports: " +  
    string.Join(", ", SerialPort.GetPortNames()));
```

Port Naming

- Windows: COM1, COM3, ...
- Linux: /dev/ttyS0, /dev/ttyUSB0, /dev/ttyACM0
- macOS: /dev/tty.usbserial-xxx,
/dev/tty.usbmodem-xxx

Basic Open/Write/Read

Minimal Console Demo

```
using System.IO.Ports;

var portName = "COM3"; // or "/dev/ttyUSB0"
using var sp = new SerialPort(portName, 115200, Parity.None,
                               8, StopBits.One) {
    Handshake = Handshake.None,
    ReadTimeout = 1000,
    WriteTimeout = 1000,
    NewLine = "\r\n" // match your device!
};
sp.Open();
Console.WriteLine($"Opened {sp.PortName} @ {sp.BaudRate} baud");
sp.WriteLine("IDN?");
var reply = sp.ReadLine(); // blocks until NewLine or timeout
Console.WriteLine($"Reply: {reply}");
```

DataReceived Event

```
using System;
using System.IO.Ports;

public class Reader {
    private readonly SerialPort _sp;

    public Reader(string name) {
        _sp = new SerialPort(name, 115200) { NewLine = "\n" };
        _sp.DataReceived += OnData;
    }
}
```

Event-Driven Reading

```
public void Start() => _sp.Open();  
public void Stop() { _sp.DataReceived -= OnData; _sp.Close(); }  
  
private void OnData(object? s,  
                    SerialDataReceivedEventArgs e) {  
    try {  
        var line = _sp.ReadExisting(); // or ReadLine() if framed  
        Console.WriteLine(line);        // avoid long work in event  
    } catch (TimeoutException) { }  
}  
}
```

Notes

- Event can fire for **partial** data; buffer and parse accordingly
- Keep handler **fast**; offload to a queue/Task if needed

Modern Pattern (no events)

```
public static class SerialAsync {  
    public static async Task RunAsync(string port,  
                                     CancellationToken ct) {  
        using var sp = new SerialPort(port, 115200);  
        sp.Open();  
        var buf = new byte[4096];  
        while (!ct.IsCancellationRequested) {  
            int n = await sp.BaseStream  
                      .ReadAsync(buf.AsMemory(0, buf.Length), ct);  
            if (n == 0) break; // closed  
            Console.Write(Encoding.ASCII.GetString(buf, 0, n));  
        }  
    }  
}
```

When to use

- High throughput
- Structured backpressure and cancellation
- Cleaner than events for streaming

NewLine and Text Encoding

- Set `SerialPort.NewLine` to your device's terminator ("`\n`", "`\r\n`", "`\r`")
- Default encoding is ASCII; change if needed:
 - `sp.Encoding = Encoding.UTF8;`
- For binary protocols: use `Read/Write(byte[], ...)` and `avoid ReadLine()`

Hex Utilities (binary debug)

```
static string Hex(byte[] a, int n) =>  
    BitConverter.ToString(a, 0, n).Replace("-", " ");
```

Handshake / RTS / DTR

```
sp.Handshake = Handshake.None;  
// or RequestToSend, XOnXOff, RequestToSendXOnXOff  
sp.RtsEnable = true;  
// manual RTS if needed  
sp.DtrEnable = true;  
// some devices require DTR high
```

Common Settings

- Baud: 9600, 19200, 38400, 57600, 115200
- Data bits: usually 8
- Parity: None/Even/Odd
- Stop bits: One/Two

Timeouts

```
sp.ReadTimeout  = 1000; // ms, throws TimeoutException  
sp.WriteTimeout = 1000;
```

- For robust protocols, implement retries and application-level checksums

Windows vs Linux/macOS

- **Windows:** install USB-serial driver if needed; COM port number via Device Manager
- **Linux:** permissions — add user to dialout (or uucp) group, or sudo
- **macOS:** use `/dev/tty.usb*=/cu.usb*`; sometimes need vendor drivers

Test Without Hardware

- Windows: **com0com** (virtual null-modem pair), Virtual Serial Driver Pro
- Linux/macOS: `socat -d -d pty,raw,echo=0 pty,raw,echo=0`

Access from One Thread

- `SerialPort` is not fully thread-safe for concurrent read/write/events
- Prefer a single I/O loop + channels/queues to communicate with UI/logic

Don't Block Event Thread

- In `DataReceived`, quickly hand off data

Proper Shutdown

```
sp.DataReceived -= OnData;  
if (sp.IsOpen) sp.Close();  
sp.Dispose();
```

- Close before Dispose; consider using `using` where possible

Common Issues

- Wrong `NewLine` → `ReadLine()` hangs
- Wrong parity/stop bits/ baud → garbage characters
- Flow control mismatch → stalled writes
- Linux permissions → “Access denied”: add user to `dialout`, re-login
- USB sleep/low-power can drop the port on laptops

System.IO.Ports

- Simple API for serial comms, event-driven or async streaming
- Mind framing (text vs binary), flow control, and timeouts
- Prefer single-threaded I/O with clean handoff to app logic
- Test with virtual ports when hardware is unavailable

Motivation

- Windows Driver Kit (WDK) = toolkit for building drivers
- Enables integration with hardware and creation of virtual devices
- Typical use cases:
 - Device drivers (USB, PCIe, Serial, Network)
 - Virtual drivers for testing
 - Filter drivers (extend behavior of existing devices)

Driver Types

- KMDF (Kernel-Mode Driver Framework)
 - Easier and safer than raw WDM
 - Standard choice for most device drivers
- UMDF (User-Mode Driver Framework)
 - Runs in user mode (safer, less powerful)
 - Good for sensors, simple USB devices
- Specialized: File system drivers, filter drivers

Workflow

- Develop driver code in Visual Studio (C, KMDF templates)
- Define device descriptors and logic (EvtDeviceAdd, queues, I/O)
- Provide an INF file → tells Windows how to install the driver
- Build → sign → install (test mode or WHQL for production)
- Debug with WinDbg/DebugView, trace with ETW

Minimal Driver (Hello WDK)

```
#include <ntddk.h>
```

```
#include <wdf.h>
```

```
DRIVER_INITIALIZE DriverEntry;
```

```
NTSTATUS DriverEntry(PDRIVER_OBJECT DriverObject,  
                   PUNICODE_STRING RegistryPath)
```

```
{  
    KdPrint(("Hello WDK! Driver loaded.\n"));  
    return STATUS_SUCCESS;  
}
```

- Logs appear in kernel debugger

[Version]

Signature="\$WINDOWS NT\$"

Class=Sample

ClassGuid={78A1C341-4539-11d3-B88D-00C04FAD5171}

Provider=%Mfg%

DriverVer=09/04/2025,1.0.0.0

[Manufacturer]

%Mfg%=Demo,NTamd64

```
[Demo.NTamd64]
%DeviceName%=Install, Root\HelloDriver
```

```
[Install]
CopyFiles=DriverCopy
```

```
[DriverCopy]
HelloDriver.sys
```

```
[Strings]
Mfg="Demo Drivers"
DeviceName="Hello WDK Driver"
```


Best Practices

- Start with KMDF templates
- Keep drivers small, move heavy logic to user mode
- Prefer to use test VMs
- Always include INF + signature (test-signed ok for dev)

Takeaways

- WDK = essential toolset for Windows drivers
- KMDF simplifies driver development
- Every driver has:
 - Code (.sys)
 - INF (install metadata)
 - Signature
- Tools: WinDbg, DebugView, TraceView
- Foundation for real hardware and virtual device projects

- LINQ: Query syntax for objects, composable operators
- Parallelism: PLINQ (parallel LINQ) and TPL (Task Parallel Library)
- Goal: Write **declarative**, **efficient** and **scalable** C# code

- Query over `IEnumerable<T>` / `IQueryable<T>`
- Deferred execution until `ToList()` / iteration
- Pure functions preferred

```
var numbers = Enumerable.Range(1, 10);  
var evensSquared = numbers  
    .Where(n => n % 2 == 0)  
    .Select(n => n * n);
```

- Filter / Map: Where, Select
- Order: OrderBy, ThenBy
- Aggregation: Count, Sum, Average
- Grouping: GroupBy, ToLookup
- Joins: Join, GroupJoin
- Quantifiers: Any, All

Group & Aggregate Example

```
var revenueByRegion = sales
    .GroupBy(s => s.Region)
    .Select(g => new {
        Region = g.Key,
        Revenue = g.Sum(s => s.Qty * s.Price)
    })
    .OrderByDescending(x => x.Revenue);
```

- Inner join with Join
- Left join with GroupJoin + DefaultIfEmpty

```
var left =  
    from c in customers  
    join o in orders on c.Id equals o.CustId into grp  
    from o in grp.DefaultIfEmpty()  
    select new { c.Name, Total = o?.Total ?? 0m };
```

- `AsParallel()` starts parallelization
- Good for CPU-bound, independent data
- Control order: `AsOrdered()` vs `AsUnordered()`
- Merge options: buffered vs unbuffered

```
var primesPar = data
    .AsParallel()
    .WithDegreeOfParallelism(Environment.ProcessorCount)
    .Where(IsPrime)
    .Count();
```


- ForAll pushes results in parallel
- WithCancellation(cts.Token) supports cancellation
- Exceptions: wrapped in AggregateException

```
data.AsParallel()  
    .Where(IsPrime)  
    .ForAll(p => bag.Add(p));
```

- `Parallel.For` / `Parallel.ForEach` for loops
- Thread-local accumulators
- Concurrency control: `ParallelOptions`

```
Parallel.For(0, arr.Length, i => arr[i] = i * i);
```

```
var opts = new ParallelOptions { MaxDegreeOfParallelism = 4 };  
Parallel.ForEach(files, opts, file => Process(file));
```

- Use Task + async/await for I/O concurrency
- Not suited for PLINQ

```
var downloads = urls.Select(async u => await http.GetStringAsync(u));  
string[] pages = await Task.WhenAll(downloads);
```

- LINQ → Declarative, sequential queries
- PLINQ → CPU-heavy, large, independent data
- TPL → Custom loops, fine-grained control
- Task/async → I/O-bound workloads

- Small data → sequential LINQ
- Heavy per-item CPU → PLINQ
- I/O → async/await Tasks
- Avoid shared state, prefer pure functions
- Measure performance!

- LINQ simplifies data queries in C#
- PLINQ = parallel LINQ, best for CPU-bound workloads
- TPL = general-purpose parallelism
- Combine wisely: pure LINQ core + parallelism at boundaries